
Continuous 3D Scene Representations with Implicit Functions

Michaël Ramamonjisoa, Van Nguyen Nguyen

ENPC's Imagine Seminars 16/12/2020

Outline

1. **Implicit functions: an illustration with 3D surface representation**
2. Neural Radiance Field (NeRF)

Explicit vs Implicit Representation (2D)

Explicit:

$$\mathbf{f}(\alpha) = (r \cos(\alpha), r \sin(\alpha))^T$$

Domain: $[0, 2\pi]$

Implicit:


$$F(x, y) = \sqrt{x^2 + y^2} - r$$

Domain: $(x, y) \in \mathbb{R}^2$

\Rightarrow Circle is implicitly defined by $\{(x, y) | F(x, y) = 0\}$

$\mathbf{f}(\alpha)$ defines the interface

$F(x, y)$ defines the **Signed Distance Function** of the circle



Explicit vs Implicit Representation (3D)

Explicit:

$$\mathbf{f}(\alpha, \beta) = (r \sin(\alpha) \cos(\beta), -r \cos(\beta), r \sin(\alpha) \sin(\beta))$$

$$\text{Domain: } \alpha \in [0; 2\pi], \beta \in [0; \pi]$$

Implicit:


$$F(x, y, z) = \sqrt{x^2 + y^2 + z^2} - r$$

$$\text{Domain: } (x, y, z) \in \mathbb{R}^3$$

\Rightarrow Sphere is implicitly defined by $\{(x, y, z) | F(x, y, z) = 0\}$

$\mathbf{f}(\alpha, \beta)$ defines the 3D surface

$F(x, y, z)$ defines the **Signed Distance Function** of the sphere



Representing 3D surfaces


Explicit:




Voxels




Point clouds



Mesh




Depth




Surface Normals

Implicit:



Signed distance field






Mixture of primitives
(e.g gaussian mixtures)

Signed Distance Field (SDF)

- Maps each 3D points \mathbf{p} to it's signed distance from to the object surface S


$$SDF(\mathbf{p}) = \min_{q \in S} \|\mathbf{p} - q\|$$

- Sign indicates whether the point \mathbf{p} is inside (-) or outside (+) of the shape
- Shape's boundary as the zero-level-set of SDF
- Allows for Constructive Solid Geometry (CSG) through boolean operations



Mixture of Gaussians

- Represents a shape as a mixture of local implicit functions (3D gaussians)



$$F(\mathbf{x}, \Theta) = \sum_{i \in [N]} f_i(\mathbf{x}, \theta_i)$$

$$f_i(\mathbf{x}, \theta_i) = c_i \exp \left(\sum_{d \in \{x, y, z\}} \frac{-(\mathbf{p}_{i,d} - \mathbf{x}_d)^2}{2\mathbf{r}_{i,d}^2} \right)$$

- Shape's boundary is defined as an iso-level of the **global** implicit function



[1] Genova19

[2] Genova20

Representing 3D surfaces with Implicit Functions

Pros:

- Compared to **point clouds**: clearly defines the (iso-)surface
- Compared to **meshes**: can continuously **adapt to arbitrary topology**
- Compared to **voxels**: can be represented with **few parameters** (e.g. mixture of simple implicit functions)
- They are **continuous** in 3D
- Can give analytic normals, can be applied with boolean operations, etc

Representing 3D surfaces with Implicit Functions


Pros:

- Compared to **point clouds**: clearly defines the (iso-)surface
- Compared to **meshes**: can continuously **adapt to arbitrary topology**
- Compared to **voxels**: can be represented with **few parameters** (e.g. mixture of simple implicit functions)
- They are **continuous** in 3D
- Can give analytic normals, can be applied with boolean operations, etc

Cons:



- SDF is well-defined for only watertight meshes (there is an interior and an exterior)
- Need extra steps to visualize

Converting Implicit Surfaces to meshes




Implicit function

Extract (zero-level) iso-surface




Mesh

Marching Cubes



Ray marching



Representing 3D surfaces with Implicit Functions

Pros:

- Compared to **point clouds**: **clearly defines the (iso-)surface**
- Compared to **meshes**: can continuously **adapt to arbitrary topology**
- Compared to **voxels**: can be represented with few parameters (e.g. mixture of simple implicit functions)
- They are **continuous** in 3D
- Can give analytic normals, can be applied with boolean operations, etc

Cons:

- Implicit functions is well-defined for only watertight meshes (there is an interior and an exterior)
- Need extra steps to visualize
- Not all complex shapes can be efficiently / accurately represented with simple primitives


Representing 3D surfaces

DeepSDF: **Efficiently** representing complex shapes by learning their SDF

Idea: Learn a **continuous** representation of 3D implicit surfaces

Query $p = (x, y, z)$, Shape latent code \mathbf{Z}

$$F(p; \mathbf{Z}) = \text{SDF}(p, \mathcal{M})$$




=> **Continuity** in 3D space **AND** shapes space

[3] Park19

Representing 3D surfaces

DeepSDF: Representing complex shapes by learning their SDF




Take home message on Implicit Functions

Representation of a continuous field

Learned implicit functions:

- Can represent complex shapes
- Are **continuous mappings** because they use **MLPs**
- Are applicable to N-D data: 2D images, 3D shapes, radiance fields




Visualization of implicit functions is done by extracting iso-surfaces:

1. Running inference for multiple queries in input space
2. Rendering the result by combining the queries

More works on Implicit Functions for 3D shape

- Occupancy Networks



[4] Mescheder 19



- PiFu and PiFuHD



[5] Saito 19

[6] Saito 20

References

- [1] Genova et al., [Learning Shape Templates with Structured Implicit Functions](#), ICCV 2019
- [2] Genova et al., [Local Deep Implicit Functions for 3D Shape](#), CVPR 2020
- [3] Park et al., [DeepSDF: Learning Continuous Signed Distance Functions for Shape Representation](#), CVPR 2019
- [4] Mescheder et al., [Occupancy Networks: Learning 3D Reconstruction in Function Space](#), CVPR 2019
- [5] Saito, Huang, Natsume et al., PIFu: [Pixel-Aligned Implicit Function for High-Resolution Clothed Human Digitization](#), ICCV 2019
- [6] Saito et al., [PIFuHD: Multi-Level Pixel-Aligned Implicit Function for High-Resolution 3D Human Digitization](#), CVPR 2020

Courses and Seminars

Lecture on [Implicit geometry](#)

Lecture on [Implicit surface](#)

Lecture on [Explicit & Implicit Surfaces](#)

[Thomas Funkhouser's talk](#) at 3DGV seminar


Princeton COS 426, Spring 2014 on [Implicit Surfaces & Solid Representations](#)

Outline

1. Implicit functions: an illustration with 3D surface representation
2. **Neural Radiance Field (NeRF)**

Scene Representation Networks: Continuous 3D-Structure-Aware Neural Scene Representations

- One of the first relevant works on scene representation, also benchmark for most of NeRF's paper

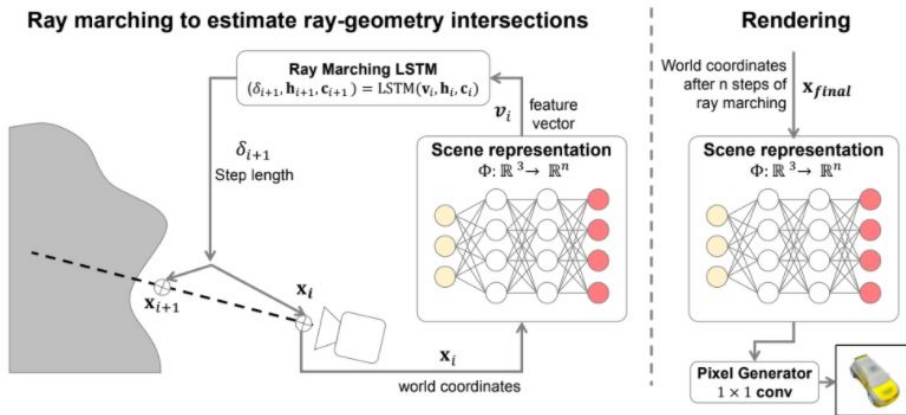


Scene Representation Networks: Continuous 3D-Structure-Aware Neural Scene Representations

- Represent a scene as a function Φ which maps a spatial location \mathbf{x} to a feature representation \mathbf{v}

$$\Phi : \mathbb{R}^3 \rightarrow \mathbb{R}^n, \quad \mathbf{x} \mapsto \Phi(\mathbf{x}) = \mathbf{v}$$

- \mathbf{v} may encode:
 - visual information: **surface color** or reflectance
 - geometry: signed distance of \mathbf{x}
- Then learn a differentiable renderer to render \mathbf{v} (using LSTM)




Definition of radiance field

- Radiance field is a 5-dimensional function which maps a 3D location \underline{x} and a direction in 3D sphere \underline{d} to a color (r, g, b) :

$$L : \mathbb{R}^3 \times S^2 \rightarrow \mathbb{R}^3$$

$$L(\underline{x}, \underline{d}) = (r, g, b)$$




- Intuitively, “radiance” is the amount of light energy passing through a given point in space, heading in a given direction
- In NeRF, there is an additional output is volume density $\sigma \in \mathbb{R}$


$$L(\underline{x}, \underline{d}) = (r, g, b, \sigma)$$

Definition of radiance field

- Radiance field is a 5-dimensional function which maps a 3D location $\mathbf{x}, \mathbf{y}, \mathbf{z}$ and a direction in 3D sphere \mathbf{d} to a color $(\mathbf{r}, \mathbf{g}, \mathbf{b})$:



(a) View 1



(b) View 2




(c) Radiance Distributions

[7] Mildenhall20

Neural Radiance field (NeRF)

Idea:

- Continuous neural networks as a view-dependent volumetric scene representation (xyz + view direction d)
- Using volumetric rendering to synthesize new views



Neural Radiance field (NeRF)

Volumetric rendering with ray tracing:

$$C(\mathbf{r}) = \int_{t_n}^{t_f} T(t) \sigma(\mathbf{r}(t)) \mathbf{c}(\mathbf{r}(t), \mathbf{d}) dt, \text{ where } T(t) = \exp\left(-\int_{t_n}^t \sigma(\mathbf{r}(s)) ds\right)$$

Opacity
Predicted colors

Volume density

Rendering model for ray $\mathbf{r}(t) = \mathbf{o} + t\mathbf{d}$:

$$C \approx \sum_{i=1}^N T_i \alpha_i c_i$$


Opacity
colors

How much light is blocked earlier along ray:

$$T_i = \prod_{j=1}^{i-1} (1 - \alpha_j)$$

How much light is contributed by ray segment i :

$$\alpha_i = 1 - e^{-\sigma_i \delta t_i}$$



[7] Mildenhall20

Neural Radiance field (NeRF)

Volumetric rendering with ray tracing:

$$C(\mathbf{r}) = \int_{t_n}^{t_f} T(t) \sigma(\mathbf{r}(t)) \mathbf{c}(\mathbf{r}(t), \mathbf{d}) dt, \text{ where } T(t) = \exp\left(-\int_{t_n}^t \sigma(\mathbf{r}(s)) ds\right)$$

Opacity
Predicted colors

Volume density

Rendering model for ray $\mathbf{r}(t) = \mathbf{o} + t\mathbf{d}$:

$$C \approx \sum_{i=1}^N T_i \alpha_i \mathbf{c}_i$$

Opacity
colors


(approximation with numerical quadrature)

How much light is blocked earlier along ray:

$$T_i = \prod_{j=1}^{i-1} (1 - \alpha_j)$$

How much light is contributed by ray segment i :

$$\alpha_i = 1 - e^{-\sigma_i \delta t_i}$$



[7] Mildenhall20

Neural Radiance field (NeRF)

Volumetric rendering with ray tracing:

$$C(\mathbf{r}) = \int_{t_n}^{t_f} T(t) \sigma(\mathbf{r}(t)) \mathbf{c}(\mathbf{r}(t), \mathbf{d}) dt, \text{ where } T(t) = \exp\left(-\int_{t_n}^t \sigma(\mathbf{r}(s)) ds\right)$$

Opacity (points to $T(t)$)
Predicted colors (points to $\mathbf{c}(\mathbf{r}(t), \mathbf{d})$)
Volume density (points to $\sigma(\mathbf{r}(t))$)

Rendering model for ray $\mathbf{r}(t) = \mathbf{o} + t\mathbf{d}$:

$$C \approx \sum_{i=1}^N T_i \alpha_i c_i$$

Opacity (points to T_i)
colors (points to c_i)

(approximation with numerical quadrature)


-> differentiable w.r.t parameters of MLP

How much light is blocked earlier along ray:

$$T_i = \prod_{j=1}^{i-1} (1 - \alpha_j)$$

How much light is contributed by ray segment i :

$$\alpha_i = 1 - e^{-\sigma_i \delta t_i}$$




[7] Mildenhall20

Neural Radiance field (NeRF)

Tricks:


- **Hierarchical Sampling**: coarse to fine importance sampling
 - First sample coarsely along the ray with stratified sampling
 - Create N_c bins between t_n and t_f
 - For each bin, sample t_i uniformly



Neural Radiance field (NeRF)

Tricks:


- **Hierarchical Sampling**: coarse to fine importance sampling
 - First sample coarsely along the ray with stratified sampling
 - Create N_c bins between t_n and t_f
 - For each bin, sample t_i uniformly



Neural Radiance field (NeRF)

Tricks:

- **Hierarchical Sampling**: coarse to fine importance sampling
 - First sample coarsely along the ray with stratified sampling
 - Create N_c bins between t_n and t_f
 - For each bin, sample t_i uniformly




(Ray-) Volume rendering

$$\hat{C}(\mathbf{r}) = \sum_{i=1}^N T_i (1 - \exp(-\sigma_i \delta_i)) \mathbf{c}_i, \text{ where } T_i = \exp\left(-\sum_{j=1}^{i-1} \sigma_j \delta_j\right)$$

Neural Radiance field (NeRF)

Tricks:

- **Hierarchical Sampling**: coarse to fine importance sampling
 - First sample coarsely along the ray with stratified sampling
 - Create N_c bins between t_n and t_f
 - For each bin, sample t_i uniformly



(Ray-) Volume rendering


$$\hat{C}(\mathbf{r}) = \sum_{i=1}^N \underbrace{T_i(1 - \exp(-\sigma_i \delta_i))}_{\mathbf{w}_i} \mathbf{c}_i, \text{ where } T_i = \exp\left(-\sum_{j=1}^{i-1} \sigma_j \delta_j\right)$$

Neural Radiance field (NeRF)

Tricks:


- **Hierarchical Sampling**: coarse to fine importance sampling

- First sample coarsely along the ray with stratified sampling
 - Create N_c bins between t_n and t_f
 - For each bin, sample t_i uniformly



(Ray-) Volume rendering $\hat{C}(\mathbf{r}) = \sum_{i=1}^N \underbrace{T_i(1 - \exp(-\sigma_i \delta_i))}_{\mathbf{W}_i} \mathbf{c}_i$, where $T_i = \exp\left(-\sum_{j=1}^{i-1} \sigma_j \delta_j\right)$

- Then do importance sampling based on color weight \mathbf{W}_i




Neural Radiance field (NeRF)

Tricks:


- **Hierarchical Sampling:** coarse to fine importance sampling

- First sample coarsely along the ray with stratified sampling
 - Create N_c bins between t_n and t_f
 - For each bin, sample t_i uniformly



(Ray-) Volume rendering $\hat{C}(\mathbf{r}) = \sum_{i=1}^N \underbrace{T_i(1 - \exp(-\sigma_i \delta_i))}_{\mathbf{w}_i} \mathbf{c}_i$, where $T_i = \exp\left(-\sum_{j=1}^{i-1} \sigma_j \delta_j\right)$

- Then do importance sampling based on color weight \mathbf{w}_i



(A) $T_i \approx 1, \sigma_i \approx 0$
 $\mathbf{w}_i \approx 0$

(B) $T_i > 0, \sigma_i > 0$
 $\mathbf{w}_i > 0$


(C) $T_i \approx 0, \mathbf{w}_i \approx 0$

Neural Radiance field (NeRF)

Tricks:


- **Hierarchical Sampling: coarse to fine importance sampling**

- First sample coarsely along the ray with stratified sampling
 - Create N_c bins between t_n and t_f
 - For each bin, sample t_i uniformly



(Ray-) Volume rendering $\hat{C}(\mathbf{r}) = \sum_{i=1}^N \underbrace{T_i(1 - \exp(-\sigma_i \delta_i))}_{\mathbf{w}_i} \mathbf{c}_i$, where $T_i = \exp\left(-\sum_{j=1}^{i-1} \sigma_j \delta_j\right)$

- Then do importance sampling based on color weight \mathbf{w}_i



(A) $T_i \approx 1, \sigma_i \approx 0$
 $\mathbf{w}_i \approx 0$

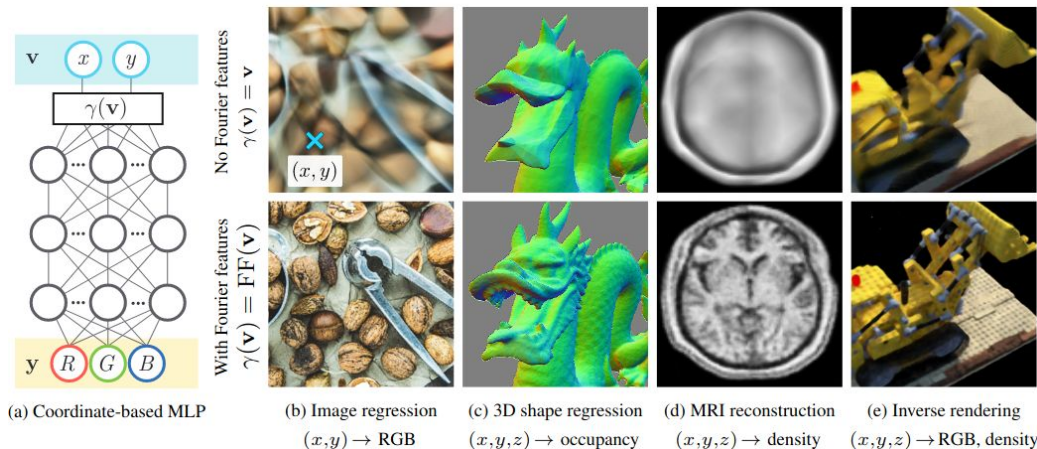
(B) $T_i > 0, \sigma_i > 0$
 $\mathbf{w}_i > 0$

(C) $T_i \approx 0, \mathbf{w}_i \approx 0$

Neural Radiance field (NeRF)

Tricks:


- **Positional encoding** to map each input 5D coordinate into a higher dimensional space
 - Learning in high-frequency mappings is difficult to learn
$$\gamma(p) = (\sin(2^0 \pi p), \cos(2^0 \pi p), \dots, \sin(2^{L-1} \pi p), \cos(2^{L-1} \pi p))$$
 - Fourier Basis feature mapping allocates neurons to different spatial frequency bands (frequency disentangling)



[8] Tancik20

Neural Radiance field (NeRF)

	Input	#Im.	L	(N_c, N_f)	PSNR \uparrow	SSIM \uparrow	LPIPS \downarrow
1) No PE, VD, H	xyz	100	-	(256, -)	26.67	0.906	0.136
2) No Pos. Encoding	$xyz\theta\phi$	100	-	(64, 128)	28.77	0.924	0.108
3) No View Dependence	xyz	100	10	(64, 128)	27.66	0.925	0.117
4) No Hierarchical	$xyz\theta\phi$	100	10	(256, -)	30.06	0.938	0.109
5) Far Fewer Images	$xyz\theta\phi$	25	10	(64, 128)	27.78	0.925	0.107
6) Fewer Images	$xyz\theta\phi$	50	10	(64, 128)	29.79	0.940	0.096
7) Fewer Frequencies	$xyz\theta\phi$	100	5	(64, 128)	30.59	0.944	0.088
8) More Frequencies	$xyz\theta\phi$	100	15	(64, 128)	30.81	0.946	0.096
9) Complete Model	$xyz\theta\phi$	100	10	(64, 128)	31.01	0.947	0.081



Ground Truth

Complete Model

No View Dependence

No Positional Encoding

Method	Diffuse Synthetic 360° [41]			Realistic Synthetic 360°			Real Forward-Facing [28]		
	PSNR \uparrow	SSIM \uparrow	LPIPS \downarrow	PSNR \uparrow	SSIM \uparrow	LPIPS \downarrow	PSNR \uparrow	SSIM \uparrow	LPIPS \downarrow
SRN [42]	33.20	0.963	0.073	22.26	0.846	0.170	22.84	0.668	0.378
NV [24]	29.62	0.929	0.099	26.05	0.893	0.160	-	-	-
LLFF [28]	34.38	0.985	0.048	24.88	0.911	0.114	24.13	0.798	0.212
Ours	40.15	0.991	0.023	31.01	0.947	0.081	26.50	0.811	0.250

[7] Mildenhall20


Neural Radiance field (NeRF)



[7] Mildenhall20

Neural Radiance field (NeRF)

NeRF in a nutshell:



- Learn the radiance field of a scene based on a collection of calibrated images
 - Use an MLP to learn continuous geometry and view-dependent appearance
- Use fully differentiable volume rendering with reconstruction loss
- Combines importance sampling and Fourier-basis encoding of 5D query to produce **high-fidelity novel view synthesis results**
- Allows efficient storage of scenes (x3000 gain over voxelized representations)

Neural Radiance field (NeRF)

Remaining challenges

- Handling dynamic scenes when acquiring calibrated views
- One network trained per scene - no generalization

Neural Radiance field (NeRF)

Remaining challenges

- Handling dynamic scenes when acquiring calibrated views
 - D-NeRF: Neural Radiance Fields for Dynamic Scenes
 - Deformable Neural Radiance Fields
- One network trained per scene - no generalization


D-NeRF: Neural Radiance Fields for Dynamic Scenes

NeRF

- Only applicable to rigid scenes
- 5D continuous function
- Requiring multiple views of a rigid scene

D-NeRF

- + Applicable for rigid and non-rigid scenes
- + 6D continuous function by considering time-component as an additional input
- + Requiring a single view per time instant for non-rigid scenes.




D-NeRF: Neural Radiance Fields for Dynamic Scenes

- **Deformation network** Ψ_t : to predict deformation field between the scene at time instant t and the scene in canonical space ($t=0$)

$$\Psi_t(\mathbf{x}, t) = \begin{cases} \Delta \mathbf{x}, & \text{if } t \neq 0 \\ 0, & \text{if } t = 0 \end{cases}$$


- **Canonical network** Ψ_x : to predict color and density in canonical configuration

$$\Psi_x(\mathbf{x}, \mathbf{d}) \mapsto (\mathbf{c}, \sigma)$$



D-NeRF: Neural Radiance Fields for Dynamic Scenes

Volumetric rendering is the same as NeRF in **canonical space**:




$$C(p, t) = \int_{h_n}^{h_f} \mathcal{T}(h, t) \sigma(\mathbf{p}(h, t)) \mathbf{c}(\mathbf{p}(h, t), \mathbf{d}) dh$$

Annotations for the equation:



- $\mathcal{T}(h, t)$ is labeled "Opacity" with an arrow pointing to it.
- $\sigma(\mathbf{p}(h, t))$ is labeled "Volume density" with an arrow pointing to it.
- $\mathbf{c}(\mathbf{p}(h, t), \mathbf{d})$ is labeled "Predicted colors" with an arrow pointing to it.

where $\mathbf{p}(h, t) = \mathbf{x}(h) + \Psi_t(\mathbf{x}(h), t)$,
 $[\mathbf{c}(\mathbf{p}(h, t), \mathbf{d}), \sigma(\mathbf{p}(h, t))] = \Psi_x(\mathbf{p}(h, t), \mathbf{d})$,
and $\mathcal{T}(h, t) = \exp \left(- \int_{h_n}^h \sigma(\mathbf{p}(s, t)) ds \right).$

D-NeRF: Neural Radiance Fields for Dynamic Scenes



Deformable Neural Radiance Fields



Deformable Neural Radiance Fields



Deformable Neural Radiance Fields vs D-NeRF

Deformable Neural Radiance Fields

Submission history

From: Keunhong Park [[view email](#)]

[v1] Wed, 25 Nov 2020 18:55:04 UTC (47,887 KB)

[v2] Thu, 26 Nov 2020 01:52:45 UTC (47,887 KB)

*We present the first method capable of photorealistically reconstructing a non-rigidly deforming scene using photos/videos captured casually from mobile phones. Our approach – **D-NeRF** – augments neural radiance fields (NeRF)*

- + Works on real data
- Relies on pretrained foreground dynamic object segmentation
- + Formulation of elastic deformation regularization
- Does not explore time dependency



D-NeRF

Submission history

From: Albert Pumarola [[view email](#)]

[v1] Fri, 27 Nov 2020 19:06:50 UTC (16,352 KB)

*ages. In this paper we introduce **D-NeRF**, a method that extends neural radiance fields to a dynamic domain, allowing to reconstruct and render novel images of objects under rigid and non-rigid motions from a single camera moving around the scene. For this purpose we consider time as an*

- Works on synthetic data
- Works on scenes with isolated object
- + Time as input

Neural Radiance field (NeRF)

Remaining challenges

- Handling dynamic scenes when acquiring calibrated views
- One network trained per scene - no generalization
 - PixelNeRF (CVPR'21 submission)
 - General radiance field (ICLR'21 submission)

PixelNeRF: Neural Radiance Fields from One or Few Images

NeRF

[7] Mildenhall20

- Optimizing NeRF of each scene independently
- Requiring many calibrated views (TODO put number estimate)
- Using canonical coordinate frame


PixelNeRF

[9] Yu20

- + Training across multiple scenes to learn a (image(s) conditioned?) scene prior
- + Address few-shot view synthesis task with sparse set of views
- + Predicting a NeRF representation in the camera coordinate system
- + **Incorporate a variable number of posed input views**


PixelNeRF: Neural Radiance Fields from One or Few Images

Incorporating multiple views:




PixelNeRF: Neural Radiance Fields from One or Few Images

Incorporating multiple views:




PixelNeRF: Neural Radiance Fields from One or Few Images

Incorporating multiple views:



PixelNeRF: Neural Radiance Fields from One or Few Images

Incorporating multiple views:



- First, transform 5D input into coordinate system of each view given camera transform


- Then, calculate intermediate feature vector for each view:

$$\mathbf{V}^{(i)} = f_1 \left(\gamma(\mathbf{x}^{(i)}), \mathbf{d}^{(i)}; \mathbf{W}^{(i)}(\pi(\mathbf{x}^{(i)})) \right)$$

The diagram shows the components of the equation: a red dot represents $\gamma(\mathbf{x}^{(i)})$, a purple arrow represents $\mathbf{d}^{(i)}$, and a blue box represents $\mathbf{W}^{(i)}(\pi(\mathbf{x}^{(i)}))$.

- Finally, aggregate with the average pooling operator ψ and passed into a the final layer

$$(\sigma, \mathbf{c}) = f_2 \left(\psi \left(\mathbf{V}^{(1)}, \dots, \mathbf{V}^{(n)} \right) \right)$$



PixelNeRF: Neural Radiance Fields from One or Few Images


	1-view			2-view		
	↑ PSNR	↑ SSIM	↓ LPIPS	↑ PSNR	↑ SSIM	↓ LPIPS
– Local	20.39	0.848	0.196	21.17	0.865	0.175
– Dirs	21.93	0.885	0.139	23.50	0.909	0.121
Full	23.43	0.911	0.104	25.95	0.939	0.071

Table 3: **Ablation studies for ShapeNet chair reconstruction.**


We show the benefit of using local features over a global code to condition the NeRF network (–Local vs Full), and of providing view directions to the network (–Dirs vs Full).

[9] Yu20

PixelNeRF: Neural Radiance Fields from One or Few Images



GRF: Learning a General Radiance Field for 3D Scene Representation and Rendering



GRF

[v1] Fri, 9 Oct 2020 14:21:43 UTC (7,696 KB)

[v2] Sun, 29 Nov 2020 06:33:25 UTC (25,183 KB)

ICLR21 submission

[OpenReview](#) grades: 7, 6, 5, 4

[I0] Trevithick20

PixelNeRF

[v1] Thu, 3 Dec 2020 18:59:54 UTC (9,768 KB)

IEEE International Conference on Neural Radiance Fields (ICNeRF)



Related Work

Lastly, note that concurrent work [42] adds image features to NeRF. A key difference is that we operate in view rather than canonical space, which makes our approach applicable in more general settings.

Moreover, we extensively demonstrate our method's performance in few-shot view synthesis, while GRF shows very limited quantitative results for this task.

More works on NeRF

- NeRF++: Analyzing and Improving Neural Radiance Fields [Zhang20]
- iNeRF: Inverting Neural Radiance Fields for Pose Estimation [Yen-Chen20]



- NeRF in the Wild [Ricardo20]...

References

- [7] Mildenhall, Srinivasan, Tancik et al., [NeRF: Representing Scenes as Neural Radiance Fields for View Synthesis](#), ECCV 2020
- [8] Tancik, Srinivasan, Mildenhall et al., [Fourier Features Let Networks Learn High Frequency Functions in Low Dimensional Domains](#), NeurIPS 2020
- [9] Yu et al., [PixelNeRF: Neural Radiance Fields from One or Few Images](#), *Arxiv preprint* 2020
- [10] Trevithick and Yang, [GRF: Learning a General Radiance Field for 3D Scene Representation and Rendering](#), *Arxiv preprint* 2020
- [11] Pumarola et al., [D-NeRF: Neural Radiance Fields for Dynamic Scenes](#), *Arxiv preprint* 2020
- [12] Park et al., [Deformable Neural Radiance Fields](#), *Arxiv preprint* 2020

- [13] Ricardo et al., [NeRF in the Wild: Neural Radiance Fields for Unconstrained Photo Collections](#), *Arxiv preprint*
- [14] Zhang et al., [NeRF++: Analyzing and improving neural radiance fields](#), *Arxiv preprint*

- [15] Yen-Chen et al., [iNeRF: Inverting Neural Radiance Fields for Pose Estimation](#), *Arxiv preprint*

Matthew Tancik's 1h [talk](#) at Tübingen seminar of the Autonomous Vision Group

Awesome Neural Radiance Fields: <https://github.com/yenchenlin/awesome-NeRF>

NeRF papers with code: <https://paperswithcode.com/method/nerf>